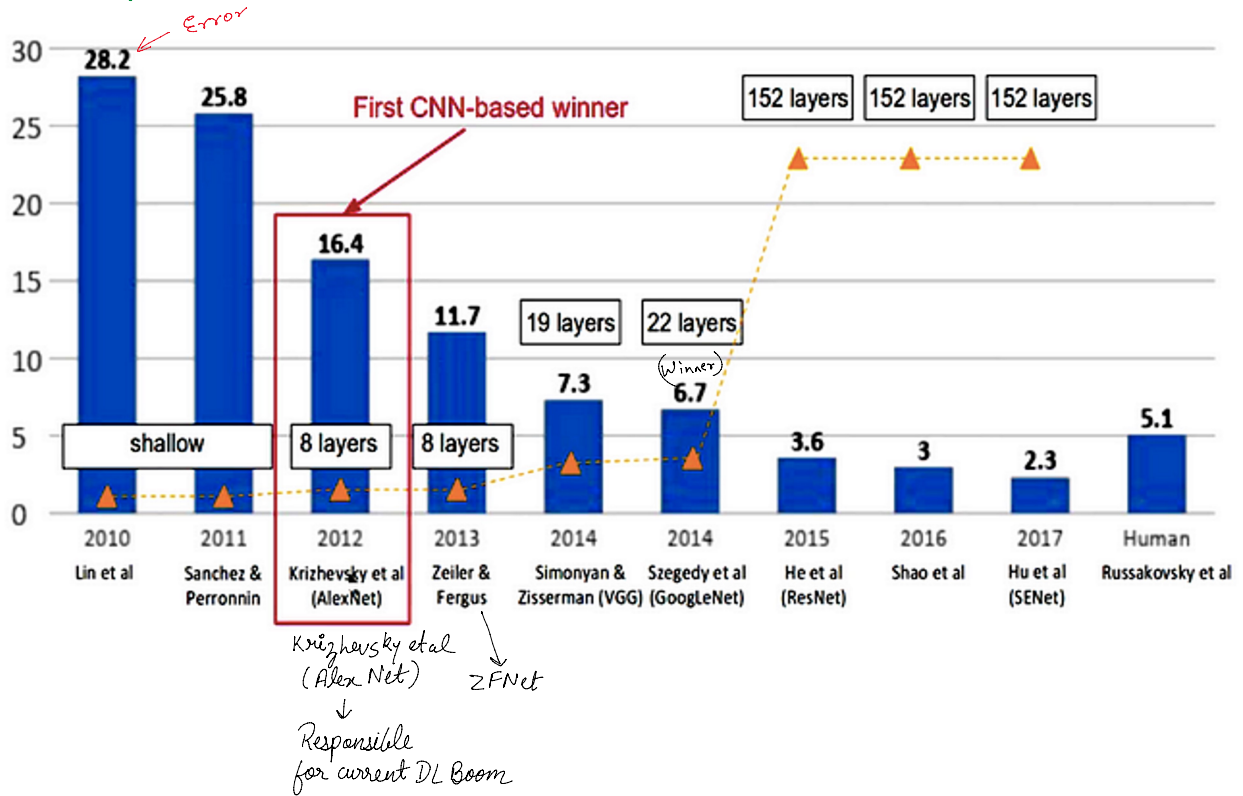




CNN Part - 3

13 September 2023 10:09 AM

Image Net challenge



Alex Net



→ Input Image Size = $224 \times 224 \times 3$

→ first layer filter = $11 \times 11 \times 3$
 stride = 4

Train set size = 1.2 M images + 60 M parameters.
 Optimization algo = mini batch G.D.

36 such filters
 ↓
 computation was divided
 across 2 cores
 ↓
 48 filters 48 filters

Trained on \rightarrow NVIDIA GTX 560 GPU (1 week to train)
 \rightarrow Activation function \rightarrow ReLU

* AlexNet introduced ReLU as an activation function.

In Alex Net, convolution layers accounted for about 90-95% computation but contained only about 5% of total parameters.

\rightarrow This is because of weight sharing, the same filter is used on the entire image \Rightarrow less parameters, but every time it is applied it is convoluted \Rightarrow more computations as it happens on entire image.

In Dense layer the parameters are more $\rightarrow (4096 \times 4096)$ Number of connections b/w first 2 Dense layers.

ZF-Net

It is similar to AlexNet, with some slight modifications

\rightarrow Low-1: changed from $(11 \times 11, 4 \text{ stride})$ to $(7 \times 7, 2 \text{ stride})$
{Filter in AlexNet}

\rightarrow Low-3,4,5: Instead of using 384, 384, 256 filters respectively, ZF-Net used 512, 1024, 512 filters respectively.

ZF-Net was made by performing really good hyperparameter tuning on Alex Net.

Google Net (won the 2014 Image Net Challenge)

Important traits of Google Net:

1. \rightarrow 22 layer architecture

2. \rightarrow No fully connected layers \rightarrow Because of a large number of parameters


3. \rightarrow Deeper Network with focus on efficiency.

\rightarrow Hence the complexity of the model is also taken into consideration, which ensures that there is no overfitting.

4. \rightarrow 5 million parameters.

Google Net was able to do all this and utilized the concept of Inception Module.


Inception Module



* Inception module is helping us to pay attention to different features of different sizes by using multiple filter/conv. layers.

* The basic idea is that rather than selecting one filter why not try all of them.

Eg



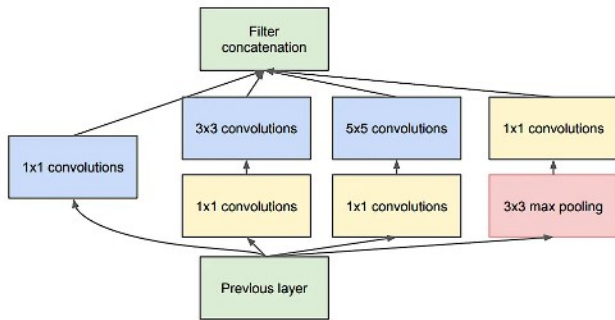


* The final results are all concatenated/stacked together.


Problems with Inception Module

↳ still computationally very expensive

Solution → Use 1×1 bottleneck layers to reduce the channel size before the expensive convolution operation.



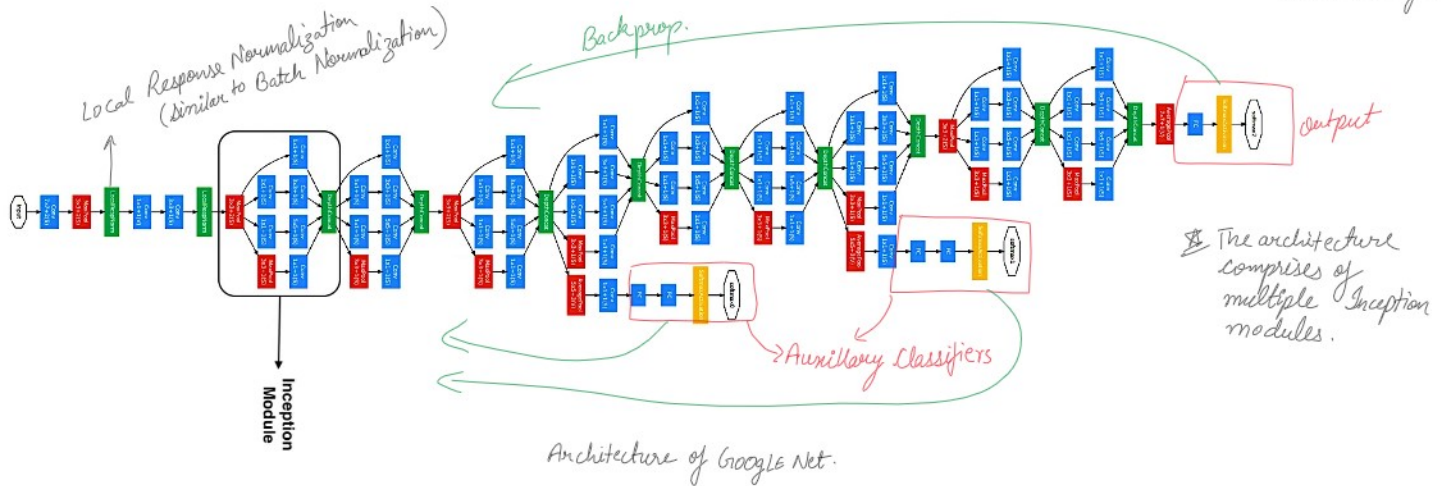
(each filter has size $1 \times 1 \times 64$, and performs a 64-dimensional dot product)



In this example we have filter of size $(1 \times 1 \times 64)$ and there are 32 such filters.

1×1 convolution layer → Bottleneck layer

⚠ Only the last output layer is fully connected layer.



⚠ The architecture comprises of multiple Inception modules.

Q What is the need for auxiliary classifiers?

Ans When performing backpropagation, since the depth of the network is too large, the gradient doesn't reach back and we run into the problem of Vanishing Gradient. In order to counter this, for the layers that are very far from the output, they not only receive gradient from the output layer but from the auxiliary classifiers as well, thereby preventing vanishing gradient. This happens by adding all the gradients received from the auxiliary classifier + the final output layer.
↳ (Ref. Backprop. in Depth)

Q In the inception module, why does the 1×1 conv. layer come after max pool layer and not before?

Ans Performing convolution is similar to getting an output pixel which is a linear combination of pixels like $w_1 x_1 + w_2 x_2 + \dots + w_d x_d$ (Refer Pic. above), now post convolution if we perform max pooling all the unique important features in the channels of input image can't be filtered out. So, by performing max pooling first we extract the important features and then by applying 1×1 conv. we can reduce the nb. of channels as in max pooling no. of channels don't change.